

Variance Reduction for Training Neural Bayes Estimators

Kai Marns-Morris

Honours (BPhil) 2026

The University of Western Australia

Supervisors: Dr Michael Bertolacci & A/Prof Edward Cripps


Background: neural Bayes estimators

Bayesian inference

- Fix a model: parameters θ with prior $p(\theta)$, and data $\mathbf{y}$ with likelihood $p(\mathbf{y} | \theta)$.
- Seeing $\mathbf{y}$, Bayes' rule updates the prior into the **posterior distribution**:


$$p(\theta | \mathbf{y}) \propto p(\mathbf{y} | \theta) p(\theta).$$

- The posterior is our full, updated belief about θ .




Point estimation

- Often we do not want the whole posterior — just a single **point estimate** $\hat{\theta}$ of θ .
- The usual choice is the **posterior mean** $\hat{\theta} = \mathbb{E}[\theta | \mathbf{y}]$.
- But it is **rarely available in closed form** — so we need a **procedure** that turns $\mathbf{y}$ into $\hat{\theta}$:




Filling the box (1): Markov chain Monte Carlo



- Accurate and general — but must be **re-run from scratch for every new y** , and can be slow.
- Do inference for *many* datasets and you pay that cost *every time*.


Filling the box (2): a neural network



- A **neural network** $\hat{\theta}_\gamma(y)$, trained *once* — a *neural Bayes estimator*.
- Training data is free: **simulate** pairs (θ, y) and adjust γ so $\hat{\theta}_\gamma(y) \approx \theta$.
- Then each new dataset is a single **forward pass** — the cost is **amortised**.

Does it work?

A simple test: estimate the mean μ and scale σ of a normal model — network vs. MCMC, on held-out data.



The network (orange) matches MCMC (blue): it has **learned the posterior-mean map**. *Why?* — the loss it was trained with.

The loss picks which summary you estimate

The network minimises the **Bayes risk** $L(\gamma) = \mathbb{E}_{\theta, \mathbf{y}}[\ell(\theta, \hat{\theta}_\gamma(\mathbf{y}))]$ — and the loss ℓ decides *which* posterior summary it targets:

loss $\ell(\theta, \hat{\theta})$	best estimate (its minimiser)
squared $\ \theta - \hat{\theta}\ ^2$	posterior mean $\mathbb{E}[\theta \mathbf{y}]$
absolute $ \theta - \hat{\theta} $	posterior median
quantile $(\alpha - \mathbf{1}_{\{\theta < \hat{\theta}\}})(\theta - \hat{\theta})$	posterior quantile

In training we **approximate** this expectation by **Monte Carlo**:

$$L(\gamma) \approx \frac{1}{N} \sum_{i=1}^N \ell(\theta_i, \hat{\theta}_\gamma(\mathbf{y}_i)), \quad (\theta_i, \mathbf{y}_i) \sim \text{model}.$$

The idea

In training we estimate the risk by **Monte Carlo**:

$$L(\gamma) \approx \frac{1}{N} \sum_{i=1}^N \ell(\boldsymbol{\theta}_i, \hat{\boldsymbol{\theta}}_{\gamma}(\mathbf{y}_i)), \quad (\boldsymbol{\theta}_i, \mathbf{y}_i) \sim \text{model}.$$

- So the **gradient** that drives training is a **random quantity with variance**.
- That variance is about *how we sample the loss*, not the inference problem — **so we can reduce it**.

What if, for part of the parameters, we could compute a conditional expectation like $\mathbb{E}[\boldsymbol{\theta}_1 \mid \boldsymbol{\theta}_2, \mathbf{y}]$ in closed form? **That is the first lever.**

Idea 1: Rao–Blackwellisation

The trick

A classical trick for killing Monte Carlo noise.

- Estimating an average by sampling? If part of it can be computed **exactly**, do that — replace the sampled piece by its exact **conditional mean**.
- Conditioning-and-averaging can only *remove* variance, never add it.

noisy draws $\longrightarrow$ replace one block by its exact mean $\longrightarrow$ tighter estimate


Named after the Rao–Blackwell theorem.

Rao–Blackwellising the loss — and why it must help


Split $\theta = (\tau, \beta)$ with β conditionally conjugate; integrate β out by replacing its sampled target with $\mathbb{E}[\beta \mid \tau, \mathbf{y}]$. For the gradient g ,

$$\text{Var}(g) = \underbrace{\text{Var}(\mathbb{E}[g \mid \tau, \mathbf{y}])}_{\text{Var}(g_{\text{RB}})} + \underbrace{\mathbb{E}[\text{Var}(g \mid \tau, \mathbf{y})]}_{\succeq 0}.$$

- Law of total variance $\Rightarrow \text{Var}(g_{\text{RB}}) \preceq \text{Var}(g)$.
- **Never increases gradient variance** — any model, any architecture.
- Linear net: the reduction reaches the fitted weights too.



Does it pay off? Bayesian linear regression



MC (blue) vs RB (orange) vs Bayes floor L^* (green) — across batch sizes 4, 16, 64, 256.

lower loss in
every grid cell

closes $\sim$ **half** (52%)
of the gap to L^*


matches the standard NBE
at $\frac{1}{4}$ **the batch size**

Idea 2: importance sampling


A more general lever

No conjugacy required — unlike Rao–Blackwell.

- Simulate from a **proposal** that oversamples the informative regions, then **reweight** back to the original risk (unbiased for any proposal).
- But **no guarantee**: bad weights can *increase* variance and shrink the effective sample size.




Importance sampling: the result



On an AR(1) model, the GRU already tracks the Gibbs (Bayes) reference across the prior — **little error left to recover.**

General, yes — but here reweighting found no error worth recovering. **No guarantee, and it did not help.**



Push the proposal hard and the effective sample size **collapses** — only adding noise.

Conclusion

Conclusion

Training a neural Bayes estimator is itself Monte Carlo — so its training variance can be reduced.

- **Rao–Blackwellisation** (needs a conjugate block): a *provable* reduction — closes $\sim$ half the gap to the Bayes floor and matches the standard estimator at $\frac{1}{4}$ the batch size.
- **Importance sampling** (needs nothing special): more general, but unguaranteed and problem-dependent — here it did not help.

Conjugacy — the classical trick behind Gibbs sampling — is just as useful for *training* the estimators meant to replace it.

Thank you